

# MPI (Beginners to Intermediate)

Monday, January 2, 2023 10:10 AM

MPI programs that you can try as a beginner or intermediate level learner:

1. **Matrix-vector multiplication:** A program that demonstrates how to perform a simple matrix-vector multiplication in parallel using MPI.
2. **Parallel sorting:** A program that demonstrates how to use MPI to sort a large dataset in parallel.
3. **Parallel prefix sum:** A program that demonstrates how to use MPI to compute the prefix sum of a large dataset in parallel.
4. **Parallel Monte Carlo integration:** A program that demonstrates how to use MPI to perform a Monte Carlo integration in parallel.
5. **Parallel dictionary:** A program that demonstrates how to use MPI to implement a simple distributed dictionary.
6. **Parallel web crawler:** A program that demonstrates how to use MPI to build a simple parallel web crawler.
7. **Parallel graph traversal:** A program that demonstrates how to use MPI to traverse a graph in parallel.
8. **Parallel image processing:** A program that demonstrates how to use MPI to perform image processing tasks in parallel.
9. **Parallel simulation:** A program that demonstrates how to use MPI to run a simple simulation in parallel.
10. **Parallel machine learning:** A program that demonstrates how to use MPI to perform machine learning tasks in parallel.

## Resources

However, there are many resources available online that can help you learn more about using MPI with Python, including tutorials, documentation, and examples.

Here are a few suggestions for finding more information about MPI and mpi4py:

- The mpi4py documentation is a good place to start learning about the library and how to use it:  
<https://mpi4py.readthedocs.io/>
- The MPI Forum website has a wealth of information about MPI and its different implementations:  
<https://www.mpi-forum.org/>
- The MPI for Python tutorial on HPC Wiki is a comprehensive guide to using MPI with Python:  
<https://hpc.gitlab.io/python/tutorials/mpi/>
- The Parallel Programming with Python tutorial on HPC Wiki has a section on using MPI with Python:  
[https://hpc.gitlab.io/python/tutorials/parallel\\_programming/](https://hpc.gitlab.io/python/tutorials/parallel_programming/)
- The High Performance Python book (O'Reilly) has a chapter on using MPI with Python:  
<https://www.oreilly.com/library/view/high-performance-python/9781449361747/ch08.html>

I hope these resources will be helpful in your journey to learn MPI with Python.

Here is the code for the "Hello world" MPI program that I provided in my previous message:

```
from mpi4py import MPI
```

```
# Get the rank of the process and the size of the communicator
```

```
rank = MPI.COMM_WORLD.rank
```

```
size = MPI.COMM_WORLD.size
```

```
# If we are the root process, send a message to the other process
```

```
if rank == 0:
```

```
    message = "Hello from process 0"
```

```
    MPI.COMM_WORLD.send(message, dest=1)
```

```
# If we are the other process, receive the message from the root process
```

```
if rank == 1:
```

```
    message = MPI.COMM_WORLD.recv(source=0)
```

```
    print(f"Received message: {message}")
```

```

from mpi4py import MPI

# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size

# If we are the root process, send a message to the other process
if rank == 0:
    message = "Hello from process 0"
    MPI.COMM_WORLD.send(message, dest=1)

# If we are the other process, receive the message from the root
process
if rank == 1:
    message = MPI.COMM_WORLD.recv(source=0)
    print(f"Received message: {message}")

```

This program sends a message from process 0 to process 1 using the `send()` function, and process 1 receives the message using the `recv()` function. The rank and size variables are used to determine which process is currently executing the code and the total number of processes in the communicator, respectively.

To run this program using `mpiexec`, you can use the following command:

**`mpiexec -n 2 python hello_world.py`**

```
mpiexec -n 2 python hello_world.py
```

**`mpiexec -n 2 python hello_world.py`**

This will start 2 processes and run the `hello_world.py` program using the `mpi4py` library. The output of the program should be:

**Received message: Hello from process 0**

```
Received message: Hello from process 0
```

## Matrix-vector multiplication:

A program that demonstrates how to perform a simple matrix-vector multiplication in parallel using MPI  
simple MPI program that demonstrates how to perform a matrix-vector multiplication in parallel:

```

from mpi4py import MPI
import numpy as np

# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size

# Create a random matrix and vector
matrix = np.random.rand(10, 10)
vector = np.random.rand(10)

# Divide the matrix into rows and distribute them among the processes
local_rows = np.array_split(matrix, size)[rank]

# Compute the local dot product
local_result = np.dot(local_rows, vector)

# Gather the local results into a single array on the root process
result = MPI.COMM_WORLD.gather(local_result, root=0)

# Print the result on the root process

```

```
if rank == 0:
    print(f"Result: {result}")
```

```
from mpi4py import MPI
import numpy as np

# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size

# Create a random matrix and vector
matrix = np.random.rand(10, 10)
vector = np.random.rand(10)

# Divide the matrix into rows and distribute them among the processes
local_rows = np.array_split(matrix, size)[rank]

# Compute the local dot product
local_result = np.dot(local_rows, vector)

# Gather the local results into a single array on the root process
result = MPI.COMM_WORLD.gather(local_result, root=0)

# Print the result on the root process
if rank == 0:
    print(f"Result: {result}")
```

This program creates a random 10x10 matrix and a random vector, and then divides the matrix into rows and distributes them among the processes using the `array_split()` function. Each process computes the dot product of its local rows and the vector using the `dot()` function, and then the results are gathered into a single array on the root process using the `gather()` function. Finally, the result is printed on the root process.

To run this program using `mpiexec`, you can use the following command:

**`mpiexec -n 4 python matrix_vector_mult.py`**

```
mpiexec -n 4 python matrix_vector_mult.py
```

**Parallel sorting:** A program that demonstrates how to use MPI to sort a large dataset in parallel

simple MPI program that demonstrates how to use MPI to sort a large dataset in parallel:

```
from mpi4py import MPI
import numpy as np

# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size

# Create a random array of size 1000
data = np.random.rand(1000)

# Sort the local data
local_data = np.sort(data[rank::size])

# Gather the local data on the root process
sorted_data = MPI.COMM_WORLD.gather(local_data, root=0)

# On the root process, concatenate and sort the gathered data
if rank == 0:
```

```
sorted_data = np.concatenate(sorted_data)
sorted_data = np.sort(sorted_data)
print(f"Sorted data: {sorted_data}")
```

```
from mpi4py import MPI
import numpy as np

# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size

# Create a random array of size 1000
data = np.random.rand(1000)

# Sort the local data
local_data = np.sort(data[rank::size])

# Gather the local data on the root process
sorted_data = MPI.COMM_WORLD.gather(local_data, root=0)

# On the root process, concatenate and sort the gathered data
if rank == 0:
    sorted_data = np.concatenate(sorted_data)
    sorted_data = np.sort(sorted_data)
    print(f"Sorted data: {sorted_data}")
```

Here is an example of a simple MPI program that demonstrates how to use MPI to sort a large dataset in parallel:

This program creates a random array of size 1000 and then sorts the data using the `sort()` function. The data is then divided among the processes using slicing, and each process sorts its local data. The local data is gathered on the root process using the `gather()` function, and

**Parallel prefix sum:** A program that demonstrates how to use MPI to compute the prefix sum of a large dataset in parallel.  
 simple MPI program that demonstrates how to use MPI to compute the prefix sum of a large dataset in parallel:

```
from mpi4py import MPI
import numpy as np

# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size

# Create a random array of size 1000
data = np.random.rand(1000)

# Compute the prefix sum of the local data
local_result = np.cumsum(data[rank::size])

# Gather the local results on the root process
results = MPI.COMM_WORLD.gather(local_result, root=0)

# On the root process, concatenate and compute the prefix sum of the gathered data
if rank == 0:
    results = np.concatenate(results)
    prefix_sum = np.cumsum(results)
    print(f"Prefix sum: {prefix_sum}")
```

```

from mpi4py import MPI
import numpy as np

# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size

# Create a random array of size 1000
data = np.random.rand(1000)

# Compute the prefix sum of the local data
local_result = np.cumsum(data[rank::size])

# Gather the local results on the root process
results = MPI.COMM_WORLD.gather(local_result, root=0)

# On the root process, concatenate and compute the prefix sum of the
gathered data
if rank == 0:
    results = np.concatenate(results)
    prefix_sum = np.cumsum(results)
    print(f"Prefix sum: {prefix_sum}")

```

This program creates a random array of size 1000 and then computes the prefix sum of the data using the `cumsum()` function. The data is then divided among the processes using slicing, and each process computes the prefix sum of its local data. The local results are gathered on the root process using the `gather()` function, and then the prefix sum of the gathered data is computed on the root process using the `cumsum()` function.

To run this program using `mpiexec`, you can use the following command:

**`mpiexec -n 4 python prefix_sum.py`**

```
mpiexec -n 4 python prefix_sum.py
```

This will start 4 processes and run the `prefix_sum.py` program using the `mpi4py` library. The output of the program should be the prefix sum of the random data array.

**Parallel Monte Carlo integration:** A program that demonstrates how to use MPI to perform a Monte Carlo integration in parallel  
 simple MPI program that demonstrates how to use MPI to perform a Monte Carlo integration in parallel:

```

from mpi4py import MPI
import numpy as np

# Define the function to be integrated
def f(x):
    return 4.0 / (1.0 + x**2)

# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size

# Set the number of samples and the integration limits
n_samples = 100000
a = 0.0
b = 1.0

# Generate a random sample of points
x = np.random.uniform(a, b, n_samples // size)

# Compute the local integral
local_result = np.sum(f(x)) * (b - a) / n_samples

# Reduce the local integrals to a single value on the root process
result = MPI.COMM_WORLD.reduce(local_result, op=MPI.SUM, root=0)

```

```
# Print the result on the root process
if rank == 0:
    print(f"Result: {result}")
```

```
from mpi4py import MPI
import numpy as np

# Define the function to be integrated
def f(x):
    return 4.0 / (1.0 + x**2)

# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size

# Set the number of samples and the integration limits
n_samples = 100000
a = 0.0
b = 1.0

# Generate a random sample of points
x = np.random.uniform(a, b, n_samples // size)

# Compute the local integral
local_result = np.sum(f(x)) * (b - a) / n_samples

# Reduce the local integrals to a single value on the root process
result = MPI.COMM_WORLD.reduce(local_result, op=MPI.SUM, root=0)

# Print the result on the root process
if rank == 0:
    print(f"Result: {result}")
```

This program defines a function  $f(x)$  to be integrated and then generates a random sample of points within the integration limits using the `uniform()` function. Each process computes the local integral using the `sum()` function and then the local integrals are reduced to a single value on the root process using the `reduce()` function. Finally, the result is printed on the root process. To run this program using `mpiexec`, you can use the following command:

**`mpiexec -n 4 python monte_carlo_integration.py`**

```
mpiexec -n 4 python monte_carlo_integration.py
```

**Parallel web crawler:** A program that demonstrates how to use MPI to build a simple parallel web crawler

MPI program that demonstrates how to use MPI to build a simple parallel web crawler:

```
from mpi4py import MPI
import requests
```

```
# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size
```

```
# Set the base URL to crawl
base_url = "http://example.com"
```

```
# Crawl the base URL
response = requests.get(base_url)
```

```
# Extract the links from the HTML content
links = extract_links(response.text)
```

```
# Divide the links among the processes
local_links = np.array_split(links, size)[rank]
```

```
# Crawl the local links
for link in local_links:
    response = requests.get(link)
```

```

# Process the response (e.g., extract data, save to database, etc.)
process_response(response)

from mpi4py import MPI
import requests

# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size

# Set the base URL to crawl
base_url = "http://example.com"

# Crawl the base URL
response = requests.get(base_url)

# Extract the links from the HTML content
links = extract_links(response.text)

# Divide the links among the processes
local_links = np.array_split(links, size)[rank]

# Crawl the local links
for link in local_links:
    response = requests.get(link)
    # Process the response (e.g., extract data, save to database,
    # etc.)
    process_response(response)

```

This program uses the requests library to crawl a given base URL and extract the links from the HTML content. The links are then divided among the processes using the `array_split()` function, and each process crawls the local links and processes the responses. To run this program using `mpiexec`, you can use the following command:

**`mpiexec -n 4 python web_crawler.py`**

```
mpiexec -n 4 python web_crawler.py
```

This will start 4 processes and run the `web_crawler.py` program using the `mpi4py` library. The processes will crawl the links in parallel and process the responses.

**Parallel dictionary:** A program that demonstrates how to use MPI to implement a simple distributed dictionary.

```
from mpi4py import MPI
```

```

# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size

```

```

# Create a dictionary on the root process
if rank == 0:
    dictionary = {"a": 1, "b": 2, "c": 3}
else:
    dictionary = None

```

```

# Scatter the dictionary keys to the processes
keys = MPI.COMM_WORLD.scatter(list(dictionary.keys()), root=0)

```

```

# On each process, print the received key and its corresponding value
print(f"Process {rank}: {keys} = {dictionary[keys]}")

```

```

from mpi4py import MPI

# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size

# Create a dictionary on the root process
if rank == 0:
    dictionary = {"a": 1, "b": 2, "c": 3}
else:
    dictionary = None

# Scatter the dictionary keys to the processes
keys = MPI.COMM_WORLD.scatter(list(dictionary.keys()), root=0)

# On each process, print the received key and its corresponding
value
print(f"Process {rank}: {keys} = {dictionary[keys]}")

```

This program creates a dictionary on the root process and then scatters the keys of the dictionary to the other processes using the `scatter()` function. Each process prints the received key and its corresponding value from the dictionary. To run this program using `mpiexec`, you can use the following command:

**`mpiexec -n 4 python distributed_dictionary.py`**

```
mpiexec -n 4 python distributed_dictionary.py
```

This will start 4 processes and run the `distributed_dictionary.py` program using the `mpi4py` library. The output of the program should be the keys and values of the dictionary printed on each process.

**Parallel graph traversal:** A program that demonstrates how to use MPI to traverse a graph in parallel

```

from mpi4py import MPI

# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size

# Create a graph on the root process
if rank == 0:
    graph = {"A": ["B", "C"], "B": ["C", "D"], "C": ["D"], "D": []}
else:
    graph = None

# Scatter the nodes of the graph to the processes
nodes = MPI.COMM_WORLD.scatter(list(graph.keys()), root=0)

# On each process, print the received node and its neighbors
print(f"Process {rank}: {nodes} -> {graph[nodes]}")

```



```

from mpi4py import MPI

# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size

# Create a graph on the root process
if rank == 0:
    graph = {"A": ["B", "C"], "B": ["C", "D"], "C": ["D"], "D": []}
else:
    graph = None

# Scatter the nodes of the graph to the processes
nodes = MPI.COMM_WORLD.scatter(list(graph.keys()), root=0)

# On each process, print the received node and its neighbors
print(f"Process {rank}: {nodes} -> {graph[nodes]}")

```

This program creates a graph on the root process and then scatters the nodes of the graph to the other processes using the `scatter()` function. Each process prints the received node and its neighbors from the graph. To run this program using `mpiexec`, you can use the following command:

**mpiexec -n 4 python graph\_traversal.py**

```
mpiexec -n 4 python graph_traversal.py
```

This will start 4 processes and run the `graph_traversal.py` program using the `mpi4py` library. The output of the program should be the nodes and their neighbors printed on each process.

**Parallel image processing:** A program that demonstrates how to use MPI to perform image processing tasks in parallel

```

from mpi4py import MPI
from PIL import Image

# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size

# Load the image on the root process
if rank == 0:
    image = Image.open("image.jpg")
else:
    image = None

# Scatter the image rows to the processes
rows = image.size[1]
local_rows = np.array_split(np.arange(rows), size)[rank]
local_image = image.crop((0, min(local_rows), image.size[0], max(local_rows)))

# Perform image processing on the local image
local_image = local_image.convert("L")
local_image = local_image.filter(ImageFilter.BLUR)

# Gather the local images on the root process
images = MPI.COMM_WORLD.gather(local_image, root=0)

# On the root process, stitch the images together
if rank == 0:
    images = [image] + images
    result = Image.new("RGB", image.size)
    y_offset = 0

```

```

for im in images:
    result.paste(im, (0, y_offset))
    y_offset += im.size[1]
result.save("result.jpg")

```

```

from mpi4py import MPI
from PIL import Image

# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size

# Load the image on the root process
if rank == 0:
    image = Image.open("image.jpg")
else:
    image = None

# Scatter the image rows to the processes
rows = image.size[1]
local_rows = np.array_split(np.arange(rows), size)[rank]
local_image = image.crop((0, min(local_rows), image.size[0], max(local_rows)))

# Perform image processing on the local image
local_image = local_image.convert("L")
local_image = local_image.filter(ImageFilter.BLUR)

# Gather the local images on the root process
images = MPI.COMM_WORLD.gather(local_image, root=0)

# On the root process, stitch the images together
if rank == 0:
    images = [image] + images
    result = Image.new("RGB", image.size)
    y_offset = 0
    for im in images:
        result.paste(im, (0, y_offset))
        y_offset += im.size[1]
    result.save("result.jpg")

```

This program uses the PIL library to load an image, divide it into rows, and distribute the rows among the processes using the `array_split()` and `crop()` functions. Each process performs image processing on its local image

**Parallel simulation:** A program that demonstrates how to use MPI to run a simple simulation in parallel

```

from mpi4py import MPI
import numpy as np

# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size

# Set the number of steps and the time step
n_steps = 1000
dt = 0.01

```

```

# Initialize the state of the system on each process
state = np.zeros(3)

# Perform the simulation
for i in range(n_steps):
    # Compute the local forces
    forces = compute_forces(state)

    # Reduce the local forces to a single value on the root process
    forces = MPI.COMM_WORLD.reduce(forces, op=MPI.SUM, root=0)

    # On the root process, update the state of the system
    if rank == 0:
        state += forces * dt

# Gather the final state on the root process
final_state = MPI.COMM_WORLD.gather(state, root=0)

# On the root process, print the final state of the system
if rank == 0:
    print(f"Final state: {final_state}")

```

```

from mpi4py import MPI
import numpy as np

# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size

# Set the number of steps and the time step
n_steps = 1000
dt = 0.01

# Initialize the state of the system on each process
state = np.zeros(3)

# Perform the simulation
for i in range(n_steps):
    # Compute the local forces
    forces = compute_forces(state)

    # Reduce the local forces to a single value on the root process
    forces = MPI.COMM_WORLD.reduce(forces, op=MPI.SUM, root=0)

    # On the root process, update the state of the system
    if rank == 0:
        state += forces * dt

# Gather the final state on the root process
final_state = MPI.COMM_WORLD.gather(state, root=0)

# On the root process, print the final state of the system
if rank == 0:
    print(f"Final state: {final_state}")

```

**Parallel machine learning:** A program that demonstrates how to use MPI to perform machine learning tasks in parallel.

```
from mpi4py import MPI
from sklearn.ensemble import RandomForestClassifier

# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size

# Load the data on the root process
if rank == 0:
    X, y = load_data()
else:
    X, y = None, None

# Scatter the data to the processes
X_local, y_local = MPI.COMM_WORLD.scatter(X, y, root=0)

# Fit a random forest classifier on the local data
clf = RandomForestClassifier()
clf.fit(X_local, y_local)

# Gather the classifiers on the root process
classifiers = MPI.COMM_WORLD.gather(clf, root=0)

# On the root process, ensemble the classifiers
if rank == 0:
    ensemble = EnsembleClassifier(classifiers)
```

```
from mpi4py import MPI
from sklearn.ensemble import RandomForestClassifier

# Get the rank of the process and the size of the communicator
rank = MPI.COMM_WORLD.rank
size = MPI.COMM_WORLD.size

# Load the data on the root process
if rank == 0:
    X, y = load_data()
else:
    X, y = None, None

# Scatter the data to the processes
X_local, y_local = MPI.COMM_WORLD.scatter(X, y, root=0)

# Fit a random forest classifier on the local data
clf = RandomForestClassifier()
clf.fit(X_local, y_local)

# Gather the classifiers on the root process
classifiers = MPI.COMM_WORLD.gather(clf, root=0)

# On the root process, ensemble the classifiers
if rank == 0:
    ensemble = EnsembleClassifier(classifiers)
```

This program uses the sklearn library to load the data, scatter it to the processes using the `scatter()` function, and fit a random forest classifier on the local data. The classifiers are then gathered on the root process using the `gather()` function, and an ensemble classifier is created on the root process by combining the classifiers.

This program defines an `EnsembleClassifier` class that combines the predictions of multiple classifiers using the `predict()` method. The data is loaded on the root process and then scattered to the processes using the `scatter()` function. Each process fits a random forest classifier on the local data, and the classifiers are gathered on the root process using the `gather()` function. Finally, the root process creates an instance of the `EnsembleClassifier` class by passing the classifiers as an argument to the constructor. To run this program using `mpiexec`, you can use the following command:

## **mpirun -n 4 python parallel\_machine\_learning.py**

```
mpirun -n 4 python parallel_machine_learning.py
```

This will start 4 processes and run the `parallel_machine_learning.py` program using the `mpi4py` library. The processes will fit random forest classifiers on the local data in parallel and then the root process will ensemble the classifiers to create a final model.